GRCResponder

Brianna Steier, Liam Gass, Elijah Tavares, Cael Howard,
June Kim, Rish Sharma, Angel Li

accenture

UCI
Donald Bren
School of Information
& Computer Sciences

Table of Contents

ICS Expo Recap

- Presented yesterday at ICS expo
- Pending assignments for the class
  - report and final demo

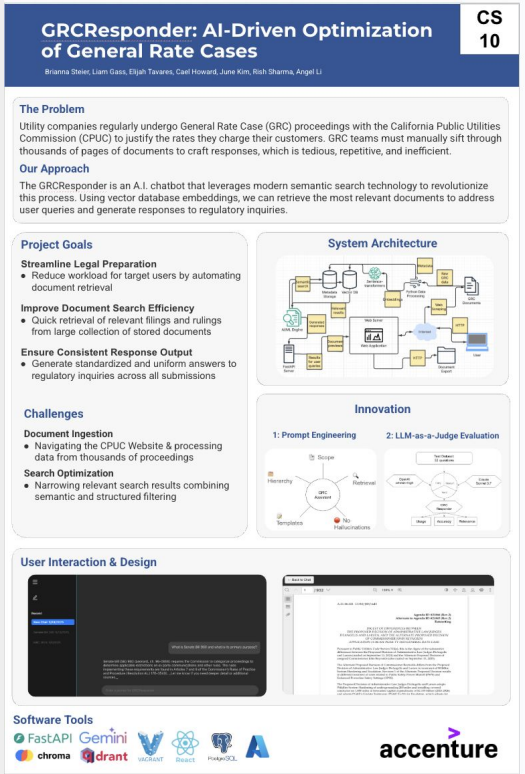

Research Updates

Current Outline

Part 4, 7, and 8 done

1. **Introduction**
   - 1.1 Context and Motivation
   - 1.2 Research Problem & Contributions
2. **Related Work**
   - 2.1 Expert Systems in Regulatory / Compliance Domains
   - 2.2 Semantic Search & RAG with LLMs
   - 2.3 Gaps GRC Addresses
3. **System Architecture**
   - 3.1 High-Level Block Diagram
   - 3.2 Data Flow
   - 3.3 Component Descriptions
4. **Implementation Details**
   - 4.1 Embedding pipeline & vector store
   - 4.2 LLM orchestration & prompt templates
   - 4.3 Back-end & front-end
   - 4.4 Performance optimizations
5. **Experimental Setup**
   - 5.1 Dataset (CPUC filings used)
   - 5.2 Baselines (keyword search)
   - 5.3 Metrics
   - 5.4 Participant details (developer lab study)
6. **Results**
   - 6.1 Quantitative results (tables/graphs)
   - 6.2 Usability findings
7. **Discussion**
   - 7.1 Interpretation & practical implications
   - 7.2 Limitations (lab vs. field)
   - 7.3 Lessons learned
8. **Conclusion**

   Recap contributions + Next steps
9. **Glossary**
10. **References**

Vector Database

Finalized Database

- Stores vectors of all documents from proceedings **filed** prior to June 25, 2025 since 2020
- Final set of documents is approximately 50 GB and 1,200 proceedings
    - Approximately 2.3 million vectors along with their text stored in the Qdrant database

Web Hosting

SSL Certificate

- unable to generate through letsencrypt due to firewall issues
- asking UCI for a new one

In the meantime…

- Switching to Docker + AWS

Questions?